

Importing the dependencies

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

```
In [5]: import os
os.getcwd()
```

```
Out[5]: 'C:\\Users\\aswin'
```

```
In [8]: os.chdir("D:\\Data Science and RPA\\Imarticus\\Dataset\\LM")
```

```
In [9]: os.getcwd()
```

```
Out[9]: 'D:\\Data Science and RPA\\Imarticus\\Dataset\\LM'
```

Importing the Dataset

```
In [19]: df=pd.read_excel("Boston.xlsx")
```

In [20]: df

Out[20]:

| | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio | black | lstat | medv |
|-----|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|--------|-------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 0.02729 | 0.0 | 7.07 | 0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | 392.83 | 4.03 | 34.7 |
| 3 | 0.03237 | 0.0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 394.63 | 2.94 | 33.4 |
| 4 | 0.06905 | 0.0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 396.90 | 5.33 | 36.2 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 501 | 0.06263 | 0.0 | 11.93 | 0 | 0.573 | 6.593 | 69.1 | 2.4786 | 1 | 273 | 21.0 | 391.99 | 9.67 | 22.4 |
| 502 | 0.04527 | 0.0 | 11.93 | 0 | 0.573 | 6.120 | 76.7 | 2.2875 | 1 | 273 | 21.0 | 396.90 | 9.08 | 20.6 |
| 503 | 0.06076 | 0.0 | 11.93 | 0 | 0.573 | 6.976 | 91.0 | 2.1675 | 1 | 273 | 21.0 | 396.90 | 5.64 | 23.9 |
| 504 | 0.10959 | 0.0 | 11.93 | 0 | 0.573 | 6.794 | 89.3 | 2.3889 | 1 | 273 | 21.0 | 393.45 | 6.48 | 22.0 |
| 505 | 0.04741 | 0.0 | 11.93 | 0 | 0.573 | 6.030 | 80.8 | 2.5050 | 1 | 273 | 21.0 | 396.90 | 7.88 | 11.9 |

506 rows × 14 columns

In [21]: *#Print first 5 rows*
df.head()

Out[21]:

| | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio | black | lstat | medv |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|--------|-------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 0.02729 | 0.0 | 7.07 | 0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | 392.83 | 4.03 | 34.7 |
| 3 | 0.03237 | 0.0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 394.63 | 2.94 | 33.4 |
| 4 | 0.06905 | 0.0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 396.90 | 5.33 | 36.2 |

```
In [23]: # Checking for number of rows and columns in our data set  
df.shape
```

```
Out[23]: (506, 14)
```

```
In [24]: # Checking for missing values in Dataset  
df.isnull().sum()
```

```
Out[24]: crim      0  
         zn        0  
         indus     0  
         chas      0  
         nox       0  
         rm        0  
         age       0  
         dis       0  
         rad       0  
         tax       0  
         ptratio   0  
         black     0  
         lstat     0  
         medv      0  
         dtype: int64
```

In [25]: *#statistical measures of the dataset*
df.describe()

Out[25]:

| | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio | |
|--------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------|
| count | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506. |
| mean | 3.613524 | 11.363636 | 11.136779 | 0.069170 | 0.554695 | 6.284634 | 68.574901 | 3.795043 | 9.549407 | 408.237154 | 18.455534 | 356. |
| std | 8.601545 | 23.322453 | 6.860353 | 0.253994 | 0.115878 | 0.702617 | 28.148861 | 2.105710 | 8.707259 | 168.537116 | 2.164946 | 91. |
| min | 0.006320 | 0.000000 | 0.460000 | 0.000000 | 0.385000 | 3.561000 | 2.900000 | 1.129600 | 1.000000 | 187.000000 | 12.600000 | 0. |
| 25% | 0.082045 | 0.000000 | 5.190000 | 0.000000 | 0.449000 | 5.885500 | 45.025000 | 2.100175 | 4.000000 | 279.000000 | 17.400000 | 375. |
| 50% | 0.256510 | 0.000000 | 9.690000 | 0.000000 | 0.538000 | 6.208500 | 77.500000 | 3.207450 | 5.000000 | 330.000000 | 19.050000 | 391. |
| 75% | 3.677083 | 12.500000 | 18.100000 | 0.000000 | 0.624000 | 6.623500 | 94.075000 | 5.188425 | 24.000000 | 666.000000 | 20.200000 | 396. |
| max | 88.976200 | 100.000000 | 27.740000 | 1.000000 | 0.871000 | 8.780000 | 100.000000 | 12.126500 | 24.000000 | 711.000000 | 22.000000 | 396. |

Understanding the correaltion between various features in Dataset

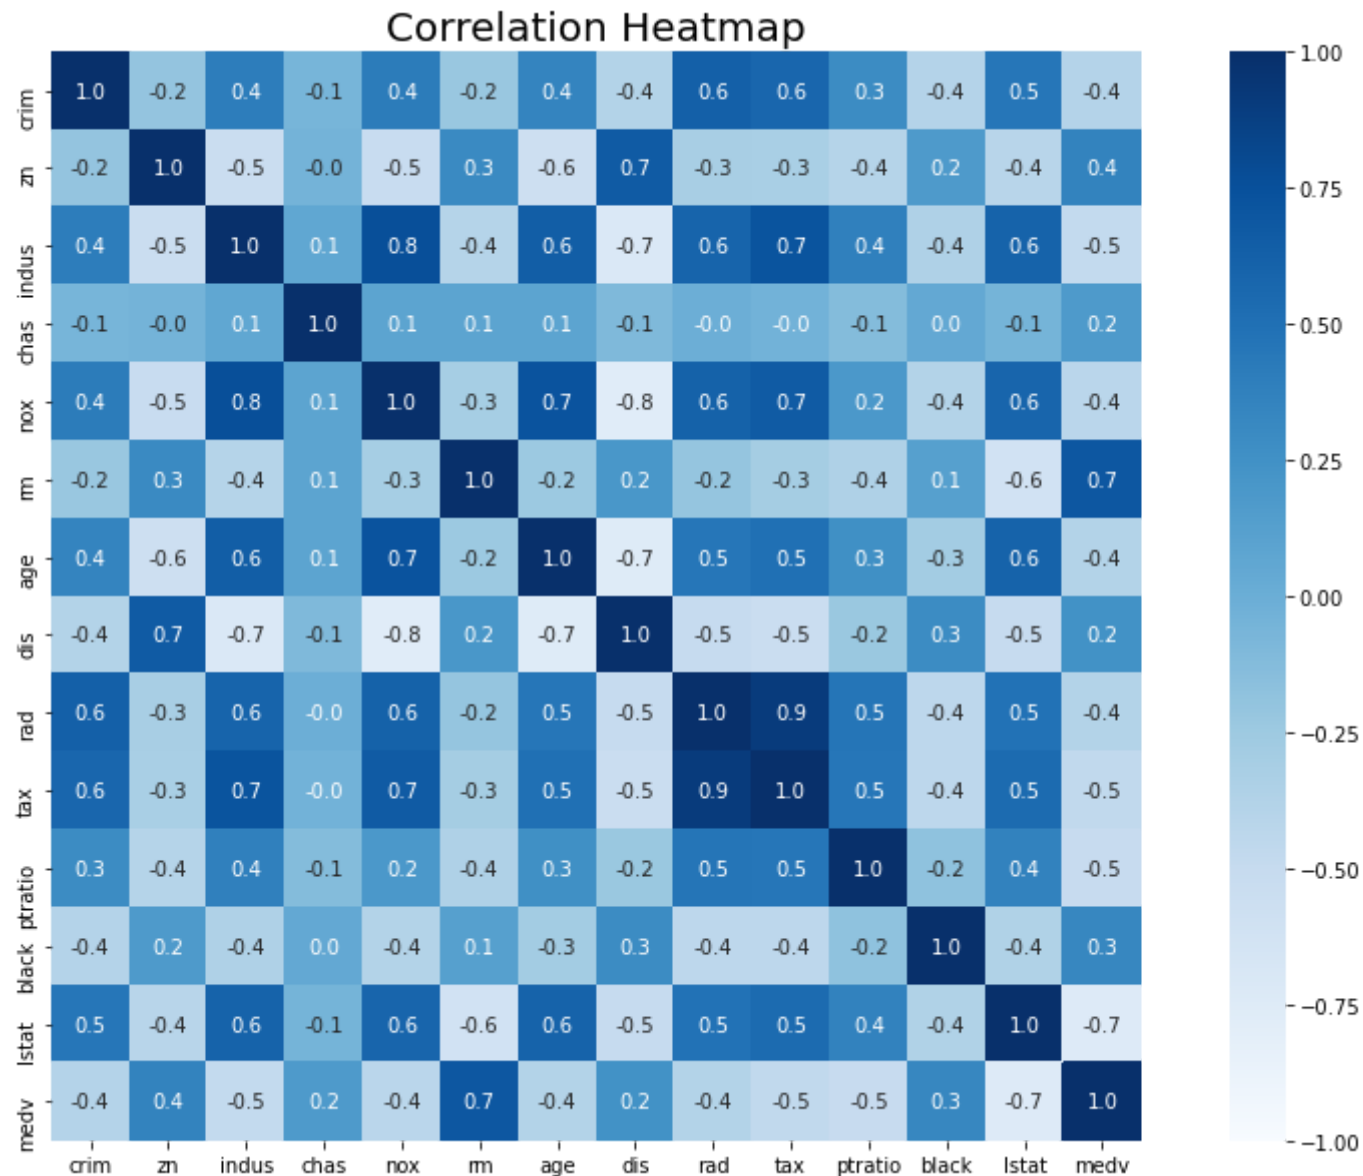
1. Positive Correlation
2. Negative Correlation

In [26]: correlation = df.corr()

In [51]: *# Constructing a heatmap to understand the correlation*

```
plt.figure(figsize=(16,10))  
sns.heatmap(df.corr(), annot=True ,fmt = '.1f' ,cbar=True,vmin=-1, vmax=1 ,square=True,cmap='Blues')  
plt.title("Correlation Heatmap",fontsize=20)
```

Out[51]: Text(0.5, 1.0, 'Correlation Heatmap')



Splitting the data and target variables

```
In [54]: x =df.drop(['medv'],axis=1)
         y =df['medv']
```

```
In [55]: x
```

```
Out[55]:
```

| | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio | black | lstat |
|------------|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|--------|-------|
| 0 | 0.00632 | 18.0 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | 396.90 | 4.98 |
| 1 | 0.02731 | 0.0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 396.90 | 9.14 |
| 2 | 0.02729 | 0.0 | 7.07 | 0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | 392.83 | 4.03 |
| 3 | 0.03237 | 0.0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 394.63 | 2.94 |
| 4 | 0.06905 | 0.0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 396.90 | 5.33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 501 | 0.06263 | 0.0 | 11.93 | 0 | 0.573 | 6.593 | 69.1 | 2.4786 | 1 | 273 | 21.0 | 391.99 | 9.67 |
| 502 | 0.04527 | 0.0 | 11.93 | 0 | 0.573 | 6.120 | 76.7 | 2.2875 | 1 | 273 | 21.0 | 396.90 | 9.08 |
| 503 | 0.06076 | 0.0 | 11.93 | 0 | 0.573 | 6.976 | 91.0 | 2.1675 | 1 | 273 | 21.0 | 396.90 | 5.64 |
| 504 | 0.10959 | 0.0 | 11.93 | 0 | 0.573 | 6.794 | 89.3 | 2.3889 | 1 | 273 | 21.0 | 393.45 | 6.48 |
| 505 | 0.04741 | 0.0 | 11.93 | 0 | 0.573 | 6.030 | 80.8 | 2.5050 | 1 | 273 | 21.0 | 396.90 | 7.88 |

506 rows × 13 columns

```
In [56]: y
```

```
Out[56]: 0      24.0  
         1      21.6  
         2      34.7  
         3      33.4  
         4      36.2  
         ...  
        501     22.4  
        502     20.6  
        503     23.9  
        504     22.0  
        505     11.9  
        Name: medv, Length: 506, dtype: float64
```

Splitting the dataset into Train and Test data

```
In [57]: X_train,X_test,Y_train,Y_test = train_test_split(x,y,test_size=0.2,random_state = 2)
```

```
In [58]: print(x.shape,X_train.shape,X_test.shape)  
  
(506, 13) (404, 13) (102, 13)
```

Model Training

XGBoost Regressor

```
In [59]: #Loading the modGBL  
         model=XGBRegressor()
```



```
In [60]: # Training the model with X_train  
model.fit(X_train,Y_train)
```

```
Out[60]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                      colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,  
                      importance_type='gain', interaction_constraints='',  
                      learning_rate=0.300000012, max_delta_step=0, max_depth=6,  
                      min_child_weight=1, missing=nan, monotone_constraints='()',  
                      n_estimators=100, n_jobs=8, num_parallel_tree=1, random_state=0,  
                      reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,  
                      tree_method='exact', validate_parameters=1, verbosity=None)
```

Evaluation of Model

Prediction on Training Data

```
In [61]: # accuracy for prediction on Training data  
training_data_prediction = model.predict(X_train)
```

```
In [62]: print(training_data_prediction)
```

| | | | | | |
|-------------|------------|------------|-----------|-----------|------------|
| [23.147501 | 20.99463 | 20.090284 | 34.69053 | 13.903663 | 13.510157 |
| 21.998634 | 15.1940975 | 10.899711 | 22.709627 | 13.832816 | 5.592794 |
| 29.810236 | 49.99096 | 34.89215 | 20.607384 | 23.351097 | 19.23555 |
| 32.695698 | 19.641418 | 26.991022 | 8.401829 | 46.00729 | 21.708961 |
| 27.062933 | 19.321356 | 19.288303 | 24.809872 | 22.61626 | 31.70493 |
| 18.542515 | 8.697379 | 17.395294 | 23.700663 | 13.304856 | 10.492197 |
| 12.688369 | 25.016556 | 19.67495 | 14.902088 | 24.193798 | 25.007143 |
| 14.900281 | 16.995798 | 15.6009035 | 12.699232 | 24.51537 | 14.999952 |
| 50.00104 | 17.525454 | 21.184624 | 31.998049 | 15.613355 | 22.89754 |
| 19.325378 | 18.717896 | 23.301125 | 37.222923 | 30.09486 | 33.102703 |
| 21.00072 | 49.999332 | 13.405827 | 5.0280113 | 16.492886 | 8.405072 |
| 28.64328 | 19.499939 | 20.586452 | 45.402164 | 39.79833 | 33.407326 |
| 19.83506 | 33.406372 | 25.271482 | 50.001534 | 12.521657 | 17.457413 |
| 18.61758 | 22.602625 | 50.002117 | 23.801117 | 23.317268 | 23.087355 |
| 41.700035 | 16.119293 | 31.620516 | 36.069206 | 7.0022025 | 20.3827 |
| 19.996452 | 11.986318 | 25.023014 | 49.970123 | 37.881588 | 23.123034 |
| 41.292133 | 17.596548 | 16.305374 | 30.034231 | 22.860699 | 19.810343 |
| 17.098848 | 18.898268 | 18.96717 | 22.606049 | 23.141363 | 33.183487 |
| 15.010934 | 11.693824 | 18.78828 | 20.80524 | 17.99983 | 19.68991 |
| 50.00332 | 17.207317 | 16.404053 | 17.520426 | 14.593481 | 33.110855 |
| 14.508482 | 43.821655 | 34.939106 | 20.381636 | 14.655634 | 8.094332 |
| 11.7662115 | 11.846876 | 18.69599 | 6.314154 | 23.983706 | 13.084503 |
| 19.603905 | 49.989143 | 22.300608 | 18.930315 | 31.197134 | 20.69645 |
| 32.21111 | 36.15102 | 14.240763 | 15.698188 | 49.99381 | 20.423601 |
| 16.184978 | 13.409128 | 50.01321 | 31.602146 | 12.271495 | 19.219482 |
| 29.794909 | 31.536846 | 22.798779 | 10.189648 | 24.08648 | 23.710463 |
| 21.991894 | 13.802495 | 28.420696 | 33.181534 | 13.105958 | 18.988266 |
| 26.576572 | 36.967175 | 30.794083 | 22.77071 | 10.201246 | 22.213818 |
| 24.483162 | 36.178806 | 23.09194 | 20.097307 | 19.470194 | 10.786644 |
| 22.671095 | 19.502405 | 20.109184 | 9.611871 | 42.799637 | 48.794792 |
| 13.097208 | 20.28583 | 24.793974 | 14.110478 | 21.701134 | 22.217012 |
| 33.003544 | 21.11041 | 25.00658 | 19.122992 | 32.398567 | 13.605098 |
| 15.1145315 | 23.088867 | 27.474783 | 19.364998 | 26.487135 | 27.499458 |
| 28.697094 | 21.21718 | 18.703201 | 26.775208 | 14.010719 | 21.692347 |
| 18.372562 | 43.11582 | 29.081839 | 20.289959 | 23.680176 | 18.308306 |
| 17.204844 | 18.320065 | 24.393475 | 26.396057 | 19.094141 | 13.3019905 |
| 22.15311 | 22.185797 | 8.516214 | 18.894428 | 21.792608 | 19.331121 |
| 18.197924 | 7.5006843 | 22.406403 | 20.004215 | 14.412416 | 22.503702 |
| 28.53306 | 21.591028 | 13.810223 | 20.497831 | 21.898977 | 23.104464 |
| 49.99585 | 16.242056 | 30.294561 | 50.001595 | 17.771557 | 19.053703 |
| 10.399217 | 20.378187 | 16.49973 | 17.183376 | 16.70228 | 19.495337 |

```

30.507633 28.98067 19.528809 23.148346 24.391027 9.521643
23.886024 49.995125 21.167099 22.597813 19.965279 13.4072275
19.948694 17.087479 12.738807 23.00453 15.222122 20.604322
26.207253 18.09243 24.090246 14.105 21.689667 20.08065
25.010437 27.874954 22.92366 18.509727 22.190847 24.004797
14.788686 19.89675 24.39812 17.796036 24.556297 31.970308
17.774675 23.356768 16.134794 13.009915 10.98219 24.28906
15.56895 35.209793 19.605724 42.301712 8.797891 24.400295
14.086652 15.408639 17.301126 22.127419 23.09363 44.79579
17.776684 31.50014 22.835577 16.888603 23.925127 12.097476
38.685944 21.388391 15.98878 23.912495 11.909485 24.960499
7.2018585 24.696215 18.201897 22.489008 23.03332 24.260433
17.101519 17.805563 13.493165 27.105328 13.311978 21.913465
20.00738 15.405392 16.595737 22.301016 24.708412 21.422579
22.878702 29.606575 21.877811 19.900253 29.605219 23.407152
13.781474 24.454706 11.897682 7.2203646 20.521074 9.725295
48.30087 25.19501 11.688618 17.404732 14.480284 28.618876
19.397131 22.468653 7.0117908 20.602013 22.970919 19.719397
23.693787 25.048244 27.977154 13.393578 14.513882 20.309145
19.306028 24.095829 14.894031 26.382381 33.298378 23.61644
24.591206 18.514652 20.900269 10.406055 23.303423 13.092017
24.675085 22.582184 20.502762 16.820635 10.220605 33.81239
18.608067 49.999187 23.775583 23.909609 21.192276 18.805798
8.502987 21.50807 23.204473 21.012218 16.611097 28.100965
21.193024 28.419638 14.294126 49.99958 30.988504 24.991066
21.433628 18.975573 28.991457 15.206939 22.817244 21.765755
19.915497 23.7961 ]

```

```

In [63]: # R-Squared Error
score_1 = metrics.r2_score(Y_train,training_data_prediction)

# Mean Absolute Error

score_2 = metrics.mean_absolute_error(Y_train,training_data_prediction)

print("R-squared error is :",score_1)
print("Mean Absolute error is :",score_2)

```


```

R-squared error is : 0.9999948236320982
Mean Absolute error is : 0.0145848437110976

```

Visualizing the actual prices and predicted prices

```
In [66]: plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.show()
```



Prediction on Test data

```
In [65]: # accuracy for prediction on Test data
test_data_prediction = model.predict(X_test)

# R-Squared Error
score_1 = metrics.r2_score(Y_test,test_data_prediction)

# Mean Absolute Error


score_2 = metrics.mean_absolute_error(Y_test,test_data_prediction)

print("R-squared error is :",score_1)
print("Mean Absolute error is :",score_2)
```

```
R-squared error is : 0.8711660369151691
Mean Absolute error is : 2.2834744154238233
```

Visualizing the Actual prices and Predicted prices for Test data

```
In [67]: plt.scatter(Y_test,test_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.show()
```



In []: